

# Intro to AI Assignment 2B

Harrison Gropper, Justin M Hu, Adideb Nag


December 14, 2022

## Question 5 - Unknown Cells


### 5c) Estimating Positions

#### Initial Probabilities



We do not know where our agent is, so there is an equal probability that it could be in each of the unblocked cells. So, our current heat-map looks like this:




For a map that looks like this:


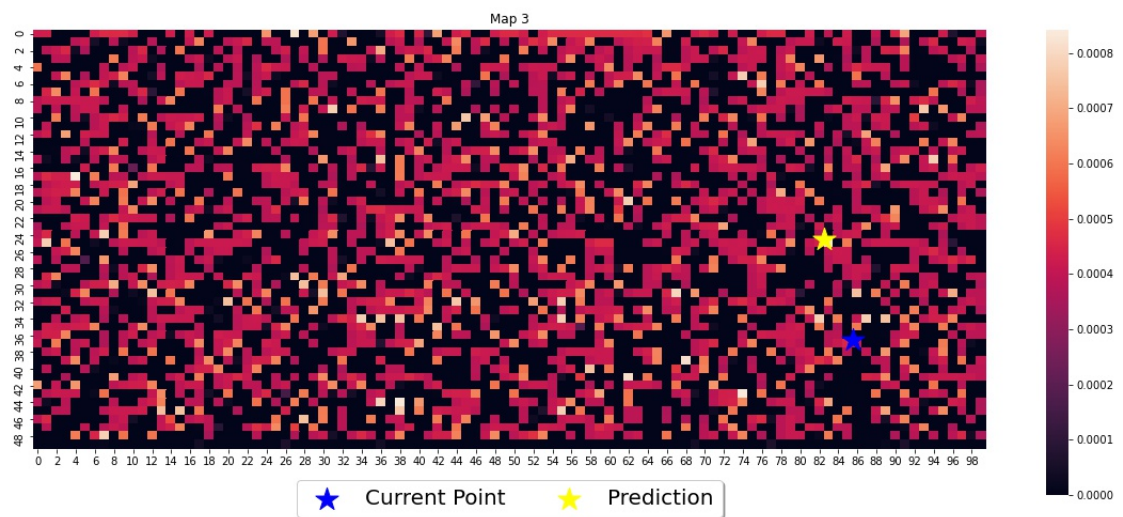


## Example Heat-maps for 10 iterations

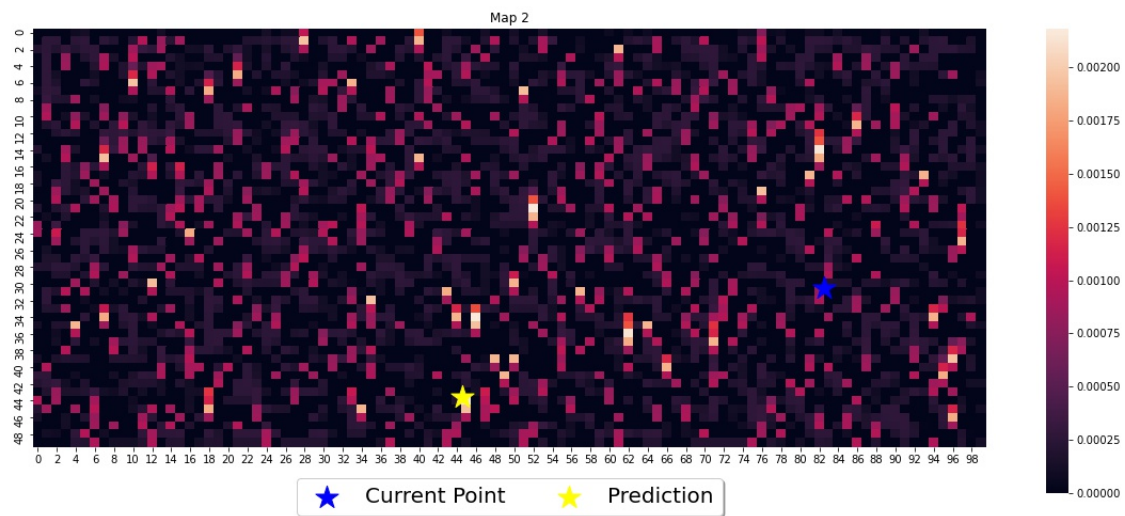




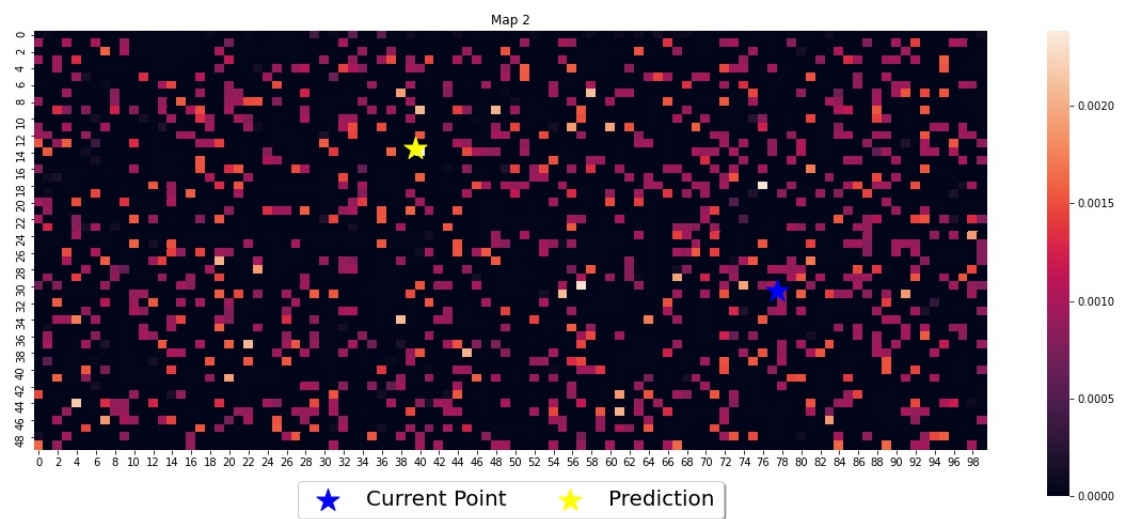

## Example Heat-maps for 50 iterations

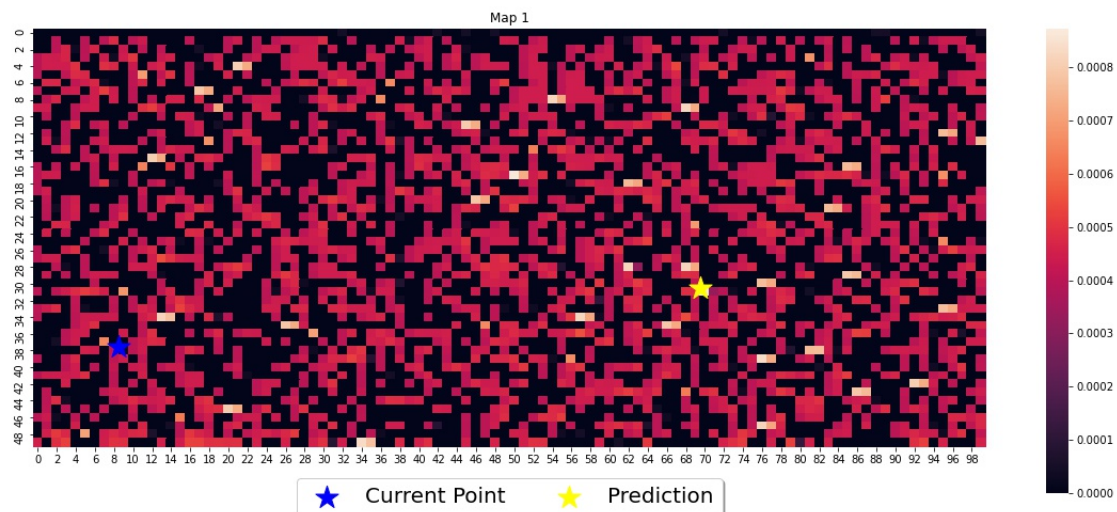




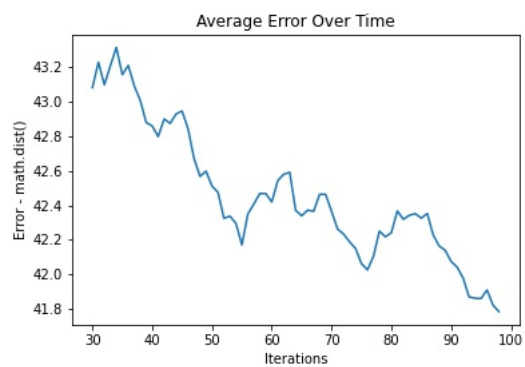


## Example Heat-maps for 100 iterations

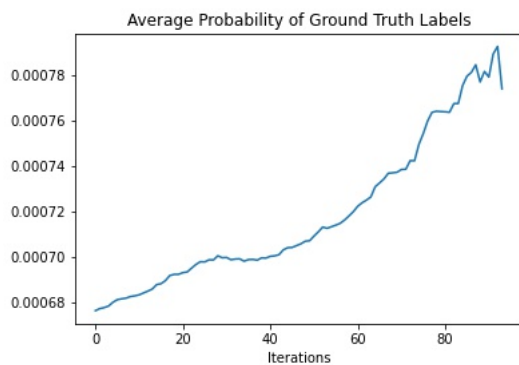




### Average Error Over All 100 Experiments



### Average Probability of Ground Truth Cells



### Extra Features

In addition to all of the computation required to make these graphs we also saved a simulation of all 100 iterations.



To see each iteration you can navigate to the animation folder and find the [GIF](#) files to view each experiment.

We have also saved each image in a folder that was used to make these GIFs.

The graph data containing the cell types and the ground truth labels were saved in a folder named data.

We have provided jupyter notebook files with well documented code, so you can see how we implemented the filtering algorithm.

We have also included a .py version of the files as well.

NOTE: The graphs used in this example were not the same graphs from assignment 2B. We generated new graphs for the 100 iterations.

## Conclusion

It seems that the less cells covered by the agent gives better results using the filtering algorithm.

When the agent is bounded (in a tight area and does not leave) the filtering algorithm works exceptionally well minimizing the overall error greatly.

This is because it is easier to rule out certain positions in the graph when our sensor reading sequences are relatively similar.

However, when our agent randomly decides to wonder off the filtering algorithm still minimizes error, but is not as accurate.

This is because we have varying sequences of evidence that confuses the filtering algorithm.

## Question 6 - The Mechanic

### 6a) Expected Net Gain From Buying $C_1$

There is a 70% chance that it is in good shape.

Then, it costs \$3000 to buy and has a \$4000 market value.

This yields a +\$1000 gain.

So, we have that  $0.70 * 1000 = 700$ .

There is a 30% chance that it is in bad shape.

We have a \$1000 in profit and we need \$1400 in repairs.

So, we have that  $0.30 * 400 = 120$ .

Our final answer is  $\$700 - \$120 = \$580$ .

The expected net gain from buying  $C_1$  given no test is \$580.

### 6b) Bayes' Theorem To Calculate Pass/Fail

Remember that  $\frac{P(B|A)P(A)}{P(B)}$ . So,

$$P(\text{pass}(c_1)) = P(\text{pass}(c_1)|q^+(c_1)) * q^+(c_1) + P(\text{pass}(c_1)|q^-(c_1)) * q^-(c_1)$$

$$P(\text{pass}(c_1)|q^+(c_1)) = 0.80 = \frac{P(q^+(c_1)|\text{pass}(c_1)) * \text{pass}(c_1)}{q^+(c_1)}$$

$$0.80 = \frac{P(q^+(c_1)|\text{pass}(c_1)) * \text{pass}(c_1)}{0.70}$$

$$P(\text{pass}(c_1)|q^-(c_1)) = 0.35 = \frac{P(q^-(c_1)|\text{pass}(c_1)) * \text{pass}(c_1)}{q^-(c_1)}$$

$$0.35 = \frac{P(q^-(c_1)|\text{pass}(c_1)) * \text{pass}(c_1)}{0.30}$$

Our answers are below:

$$P(q^+(c_1)|pass(c_1)) = \frac{P(pass(c_1)|q^+(c_1))P(q^+(c_1))}{P(pass(c_1))}$$

$$= \frac{0.80*0.70}{0.80*0.70+0.35*0.30} = 0.84$$

$$P(q^-(c_1)|pass(c_1)) = \frac{P(pass(c_1)|q^-(c_1))P(q^-(c_1))}{P(pass(c_1))}$$

$$= \frac{0.35*0.30}{0.80*0.70+0.35*0.30} = 0.158$$

$$P(q^+(c_1)|\neg pass(c_1)) = \frac{P(\neg pass(c_1)|q^+(c_1))P(q^+(c_1))}{P(\neg pass(c_1))}$$

$$= \frac{(1-0.80)*(0.70)}{1-0.80*0.70+0.35*0.30} = 0.418$$

$$P(q^-(c_1)|\neg pass(c_1)) = \frac{P(\neg pass(c_1)|q^-(c_1))P(q^-(c_1))}{P(\neg pass(c_1))}$$

$$= \frac{(1-0.35)*0.30}{1-(0.80*0.70+0.35*0.30)} = 0.58$$

## 6c) Best Decision and Expected Utility

Given that the car passes the mechanic's test the best decision is to buy the car since there is an 84% chance that it is in good condition and since the expected utility is positive

Expected Utility:

$$0.84 * (\$4000 - \$3000) + 0.16(\$4000 - \$3000 - \$1400) - \$100$$

$$= \$840 - \$64 - \$100$$

$$= \$676$$

Given that the car fails the mechanic's test, the best decision is to buy the car because even though it is only 41.80% that the car is in good shape, the expected utility is still positive

Expected Utility:

$$0.418 * (\$4000 - \$3000) + 0.582(\$4000 - \$3000 - \$1400) - \$100$$

$$= \$418 - \$232.80 - \$100$$

$$= \$85.20$$

## 6d) The Value of Optimal Information for the Mechanic's Test

Value of Optimal Information = Utility with Mechanic - Utility without Mechanic

Value of Optimal Information = \$676 - \$580 = \$96

It is worth going to the mechanic because the value of optimal information is positive even with the \$100 paid for the test accounted for in the utility equations.

## Question 7 - Markov Decision Process

### Original Utilities

In the code, we use the given initial values to iterate through the Value Iteration Formula. Using this process, we created a PossibleActions class, each holding two action names and corresponding weights. It continues doing this, updating the information until the difference goes past a given *max\_error* value.

### Five Intermediate Results

Our code starts with the given default utility values:

$$V(s_1) = 2.0$$

$$V(s_2) = 1.0$$

$$V(s_3) = 2.0$$

$$V(s_4) = 2.0$$

And the given values for the other variables:

$max\_error = 0.001$

$discountFactor = 0.75$

The following are 5 Iterations of the Code:

Iteration 6 :

$u_1 : 1.1650237910156256$

$u_2 : 1.5587219203125005$

$u_3 : 2.243067409667969$

$u_4 : 1.6079361086425783$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

Iteration 7 :

$u_1 : 1.1395152338378909$

$u_2 : 1.5796487338476566$

$u_3 : 2.205952081481934$

$u_4 : 1.6346657096740724$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

Iteration 8 :

$u_1 : 1.1517267108801272$

$u_2 : 1.5605185589663093$

$u_3 : 2.2259992822555543$

$u_4 : 1.611617583225861$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

Iteration 9 :

$u_1 : 1.1397295565675358$

$u_2 : 1.569677353198279$

$u_3 : 2.208713187419396$

$u_4 : 1.623420834264439$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

Iteration 10 :

$u_1 : 1.1450119340437939$

$u_2 : 1.5606795154313797$

$u_3 : 2.217565625698329$

$u_4 : 1.6126379640779254$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

Iteration 11:

$u_1 : 1.1393345685533245$

$u_2 : 1.5646413027337045$

$u_3 : 2.209478473058444$

$u_4 : 1.6178046446522167$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

## Final Results

The Last Iteration was: Iteration 17:

$u_1 : 1.1382596177940973$

$u_2 : 1.5600493275743013$

$u_3 : 2.2092827785843383$

$u_4 : 1.612773041296828$

$Policy : s_1 - 2, s_2 - 2, s_3 - 3, s_4 - 1$

and printed at the end:

Iterations: 18

Runtime: 12910 microseconds

## Computation Time and Number of Iterations

As shown, with the setup that the code follows, it takes 18 iterations and 12910 microseconds to compute the following output.

## Code Implementation

The Code looks as follows:

The Question7 Class does most of the work and looks like this:

```
1 package Project2;
2
3 public class Question7 {
4     public static PossibleAction[] actions = new PossibleAction[4]; // List of Actions that can be taken and their
5                                     // properties
6     public static double[] utils = new double[4]; // Utility Values
7     public static double[] utils_next = new double[4]; // Next utility values
8     public static double[] reward = new double[4]; // Calculated reward function output
9     public static char[] policy = new char[4]; // Policy character array
10
11     // Populates the graph based on the diagram provided in the problem
12     public static void populate() {
13         utils[0] = 2.0;
14         utils[1] = 1.0;
15         utils[2] = 2.0;
16         utils[3] = 2.0;
17         reward[0] = 0.0;
18         reward[1] = 0.0;
19         reward[2] = 1.0;
20         reward[3] = 0.0;
21
22         PossibleAction action0 = new PossibleAction(); // creates new PossibleAction Object and adds it to the
23         action0.action_name1 = '1';
24         action0.weight1[0] = 0.2;
25         action0.weight1[1] = 0.8;
26         action0.weight1[2] = 0.0;
27         action0.weight1[3] = 0.0;
28         action0.action_name2 = '2';
29         action0.weight2[0] = 0.2;
30         action0.weight2[1] = 0.0;
31         action0.weight2[2] = 0.0;
32         action0.weight2[3] = 0.8;
33         actions[0] = action0;
34     }
```

```

34
35     PossibleAction action1 = new PossibleAction();
36     action1.action_name1 = '2';
37     action1.weight1[0] = 0.0;
38     action1.weight1[1] = 0.2;
39     action1.weight1[2] = 0.8;
40     action1.weight1[3] = 0.0;
41     action1.action_name2 = '3';
42     action1.weight2[0] = 0.8;
43     action1.weight2[1] = 0.2;
44     action1.weight2[2] = 0.0;
45     action1.weight2[3] = 0.0;
46     actions[1] = action1;
47
48     PossibleAction action2 = new PossibleAction();
49     action2.action_name1 = '4';
50     action2.weight1[0] = 0.0;
51     action2.weight1[1] = 1.0;
52     action2.weight1[2] = 0.0;
53     action2.weight1[3] = 0.0;
54     action2.action_name2 = '3';
55     action2.weight2[0] = 0.0;
56     action2.weight2[1] = 0.0;
57     action2.weight2[2] = 0.0;
58     action2.weight2[3] = 1.0;
59     actions[2] = action2;
60
61     PossibleAction action3 = new PossibleAction();
62     action3.action_name1 = '1';
63     action3.weight1[0] = 0.0;
64     action3.weight1[1] = 0.0;
65     action3.weight1[2] = 0.9;
66     action3.weight1[3] = 0.1;
67     action3.action_name2 = '4';
68     action3.weight2[0] = 0.8;
69     action3.weight2[1] = 0.0;
70     action3.weight2[2] = 0.0;
71     action3.weight2[3] = 0.2;
72     actions[3] = action3;
73 }
74
75 // Calculates the upcoming utility values
76 public static void nextUtil(double discount_factor) {
77     double util1;
78     double util2;
79
80     // calculating the new util1 and util2 values for each i value
81     for (int i = 0; i < 4; i++) {
82         util1 = reward[i];
83         for (int j = 0; j < 4; j++) {
84             util1 += discount_factor * actions[i].weight1[j] * utils[j];
85         }
86
87         util2 = reward[i];
88         for (int j = 0; j < 4; j++) {
89             util2 += discount_factor * actions[i].weight2[j] * utils[j];
90         }
91
92         if (util1 >= util2) { //comparing the util1 and util2 value for new policy value
93             utils_next[i] = util1;
94             policy[i] = actions[i].action_name1;
95         } else {
96             utils_next[i] = util2;
97             policy[i] = actions[i].action_name2;
98         }
99     }
100 }
101

```



```

101
102 public static void main(String[] args) {
103     if (args.length < 2) { //checking if values for both max_error and discount_factor are given within arguments
104         System.out.println("Not enough arguments");
105         return;
106     } else {
107         double max_error = Double.parseDouble(args[0]);
108         double discount_factor = Double.parseDouble(args[1]);
109         if (max_error < 0.0) {
110             System.out.println("Error");
111             return;
112         }
113         long start = System.nanoTime(); //keeping track of the time (start value)
114
115         populate(); //populating the graph with the necessary information
116         int iterations = 0; //counting the number of iterations
117
118         while (true) {
119             if (iterations != 0)
120                 System.out.println("\n");
121
122             System.out.println("Iteration " + iterations + ": ");
123             nextUtil(discount_factor);
124
125             System.out.println("u1: " + utils_next[0] + "\nu2: " + utils_next[1] + "\nu3: " + utils_next[2]
126                 + "\nu4: " + utils_next[3]); //printing values of updated util values
127             System.out.println("Policy: s1 -> " + policy[0] + ", s2 -> " + policy[1] + ", s3 -> " + policy[2]
128                 + ", s4 -> " + policy[3]); //printing the policy for this iteration
129
130             iterations++;
131
132             double max_diff = 0.0; //calculating the difference between the util values
133             for (int i = 0; i < 4; i++) {
134                 double diff = utils[i] - utils_next[i];
135                 if (diff < 0) {
136                     diff *= -1;
137                 }
138                 if (diff > max_diff) {
139                     max_diff = diff;
140                 }
141                 utils[i] = utils_next[i];
142             }
143
144             if (max_diff <= max_error) { //checking if the difference falls below the max_error
145                 break;
146             }
147         }
148         long end = System.nanoTime(); //stopping time for timer
149         int elapsed = (int)(end - start) / 1000; //calculating processing time in microseconds
150
151         System.out.println("\n\nIterations: " + iterations);
152         System.out.println("Runtime: " + elapsed + " microseconds");
153     }
154 }
155 }
156 }

```

And the PossibleActions Class looked like this:

```

1 package Project2;
2
3
4 public class PossibleAction
5 {
6     char action_name1; //first action name
7     double[] weight1; //first action weight
8     char action_name2; //second action name
9     double[] weight2; //second action weight
10
11     public PossibleAction() { //Constructor to create the object and initialize all values
12         action_name1 = ' ';
13         weight1 = new double[4];
14         action_name2 = ' ';
15         weight2 = new double[4];
16     }
17 }
18

```

The code implemented for question 7 has been provided in our submission.